RAJALAKSHMI
ENGINEERING COLLEGE
An AUTONOMOUS Institution
Affiliated to ANNA UNIVERSITY, Chennai

E-COMMERCE CHURN PREDICTION

Submitted by:

DEVIKA C (221801009)
KEERTHIKA P (221801027)
LAVANYA S (221801028)

AD19541 SOFTWARE ENGINEERING METHODOLOGY

Department of Artificial Intelligence and Data Science

Rajalakshmi Engineering College, Thandalam

BONAFIDE CERTIFICATE

Certified that this project report "**E-COMMERCE CHURN PREDICTION"** is the bonafide work of "**DEVIKA C (221801009),KEERTHIKA P (221801027), LAVANYA S(221801028)**"
who carried out the project work under my supervision.

Submitted for the Practical Examination held on __________________

SIGNATURE	**SIGNATURE**
Dr.J.M.GNANASEKAR **Professor & Head,** **Dept. of Artificial Intelligence and Data Science,** **Rajalakshmi Engineering College,** **Thandalam, Chennai-602105**	**Dr.MANORANJINI J** **Associate Professor,** **Dept. of Artificial Intelligence and Data Science ,** **Rajalakshmi Engineering College,** **Thandalam, Chennai-602105**

INTERNAL EXAMINER **EXTERNAL EXAMINER**

ABSTRACT

Customer churn prediction is critical in the e-commerce industry, where understanding and retaining customers are essential to maintaining growth and profitability. To address this, a predictive model using the CatBoost algorithm—a gradient boosting model optimized for categorical data—was developed to identify customers at risk of disengaging. By analyzing a comprehensive set of customer data, including purchase history, browsing patterns, engagement levels, and transaction behaviors, the model identifies significant factors that contribute to churn, offering insights into customer behavior and loyalty risks.

A key feature of this work is an interactive dashboard that visualizes churn predictions, customer segments, and trends, making the model's insights accessible and actionable for business stakeholders. This dashboard allows companies to monitor churn rates, segment customers by risk, and observe behavioral trends over time. Equipped with these insights, e-commerce businesses can tailor retention strategies to target at-risk customers, such as offering personalized incentives or improving engagement efforts. This approach not only enhances churn prediction accuracy but also supports data-driven customer relationship management, ultimately fostering sustainable growth by minimizing revenue loss from customer attrition.

LIST OF FIGURES

1. INTRODUCTION

1.1.GENERAL OVERVIEW:

This project focuses on using the CatBoost algorithbm, a gradient boosting model known for its efficiency in handling categorical data, to predict customer churn in an e-commerce context. By analyzing various customer behaviors, like purchase history, browsing patterns, and engagement levels, the model identifies key factors associated with churn. This predictive power is complemented by a comprehensive dashboard that provides visual insights, helping e-commerce businesses make data-driven retention decisions.

1.2.NEED FOR THE SURVEY:

Understanding customer churn is critical for sustaining growth and profitability in e-commerce. Churn prediction helps identify disengaged customers early, allowing businesses to implement targeted strategies to retain them. The survey serves to gather insights into the effectiveness of these predictive models and to assess customer engagement, satisfaction, and reasons for discontinuation, supporting the development of enhanced retention tactics.

1.3.OVERVIEW OF THE PROJECT:

The main goal of this project is to build a powerful churn prediction model using the CatBoost algorithm, known for its accuracy and effectiveness with categorical data. This model will analyze various customer data points, such as purchase history, browsing patterns, and engagement metrics, to identify customers who are at risk of discontinuing their purchasing activity. Additionally, a comprehensive dashboard will be developed to visualize these insights, displaying churn metrics, customer segments, and predictive trends. This dashboard will provide e-commerce businesses with a user-friendly tool to understand churn patterns and make data-driven decisions to improve customer retention.

1.4. OBJECTIVES OF THE SURVEY:

The survey is designed to accomplish several important goals that will aid in the development, evaluation, and refinement of a customer churn prediction system. Each objective addresses a different aspect of understanding and mitigating customer churn. The survey's results will be instrumental in refining the CatBoost-based prediction model and guiding business strategies for customer retention. The specific objectives of the survey are as follows:

1. **Identify Key Factors Influencing Churn:** One of the primary goals of this survey is to identify the key factors that influence customer churn. By gathering responses on customer behaviors, preferences, and transactional patterns, the survey aims to pinpoint specific metrics that signal a higher likelihood of customers leaving. This could include metrics such as order frequency, usage patterns, and interaction with customer service. Understanding these factors is essential for fine-tuning the CatBoost model so that it can accurately assess churn risk based on real customer behaviors and trends. With this information, businesses can prioritize the most critical indicators of churn and proactively address potential issues.

2. **Evaluate Model Accuracy:** Another objective is to gauge the accuracy and effectiveness of the CatBoost-based churn prediction model. By collecting feedback from stakeholders, including analysts and decision-makers, the survey will provide insights into how well the model performs in predicting churn. Accuracy is vital to ensure that the model is a reliable tool for business decisions; inaccurate predictions could lead to misallocated resources or missed opportunities to retain valuable customers. The survey results will help determine if the model meets the organization's standards and identify areas for improvement, such as recalibrating the algorithm or adjusting parameters to increase predictive accuracy.

3. **Refine Retention Strategies:** The survey aims to gather actionable insights that can help businesses tailor their retention strategies more effectively. Predictive information alone is not enough; it must be aligned with targeted strategies to prevent churn. By understanding which factors are most influential in predicting churn, businesses can develop specific interventions, such as loyalty programs, personalized offers, or improved

customer support, to retain at-risk customers. The survey will provide valuable feedback on which strategies may resonate best with customers, enabling businesses to optimize their approach to retention and ensure that interventions are both effective and cost-efficient.

4. **Gain Customer Insights:** Finally, the survey will serve as a tool to gain deeper insights into customer satisfaction, expectations, and pain points. By analyzing customer feedback, businesses can identify common issues or unmet needs that might contribute to churn. For instance, customers might express dissatisfaction with product quality, service responsiveness, or pricing structures. These insights will not only inform the predictive model but also highlight areas where the business can improve its overall offering. Addressing these issues proactively can lead to a more positive customer experience, potentially reducing churn rates over time. Additionally, understanding customer expectations can lead to continuous improvements in both the predictive model and the retention strategies, creating a more customer-centric approach to churn prevention.

CHAPTER 2

2. REVIEW OF THE LITERATURE

2.1 INTRODUCTION

In the rapidly evolving digital marketplace, customer retention has become a vital aspect of sustaining profitability for e-commerce businesses. As competition grows, acquiring new customers is increasingly costly, making it essential for companies to focus on retaining their existing customer base. Customer churn—defined as the rate at which customers cease their relationship with a company—has a significant impact on an e-commerce business's revenue, growth, and overall market presence. Thus, identifying customers at risk of churning has become an essential goal for e-commerce platforms striving to maximize customer lifetime value. Accurately predicting which customers are likely to churn enables companies to take timely and targeted actions that can reduce churn rates, improve customer satisfaction, and ultimately enhance customer loyalty.

Churn prediction is complex, involving the analysis of diverse customer behaviors and preferences. E-commerce platforms capture a vast array of data, including transactional history, browsing patterns, and engagement metrics. This data provides valuable insights into customer behavior and helps businesses understand the factors that contribute to customer loyalty or churn. Yet, traditional approaches to addressing churn, such as general promotions or loyalty programs, often fall short of expectations. These approaches may not reach the specific customers who are likely to leave and can even result in unnecessary resource allocation on customers who have a low churn risk. As a result, modern data-driven approaches, specifically using machine learning algorithms, have gained popularity for predicting churn in a more personalized and effective way.

This project proposes the use of the CatBoost algorithm, a state-of-the-art gradient boosting algorithm, to build a high-accuracy churn prediction model for e-commerce. Gradient boosting algorithms are particularly well-suited to classification tasks, where the objective is to categorize customers as high or low risk of churning. CatBoost has been chosen for its ability to handle large datasets, work efficiently with categorical features (which are common in e-commerce data), and provide strong predictive performance with minimal hyperparameter tuning. These features make CatBoost an ideal choice for a churn prediction model that requires both accuracy and scalability.

The primary goal of this project is to create a predictive model that can accurately assess each customer's likelihood of churning based on key behavioral and engagement metrics. By analyzing past customer data—such as the frequency and recency of purchases, the types of products viewed and purchased, and the extent of interaction with marketing emails—the model will be able to predict the likelihood that a customer will discontinue their engagement with the platform. Furthermore, this project seeks to understand the underlying factors that drive churn, allowing e-commerce businesses to gain a deeper insight into customer behavior and make more informed decisions.

In addition to developing an accurate model, this project emphasizes the importance of making the results accessible and actionable through a comprehensive dashboard. This dashboard will present churn predictions in a visual format, highlighting at-risk customers and providing insights into the primary factors influencing churn for each customer segment. The goal of the dashboard is to offer a user-friendly tool that allows business stakeholders, including marketing and customer service teams, to quickly identify trends, understand churn drivers, and implement targeted retention strategies.

The predictive model and accompanying dashboard aim to empower e-commerce businesses with insights that can inform proactive customer retention strategies. Rather than relying on reactive measures, businesses will be able to intervene early, addressing issues that may lead to churn before customers decide to leave. For example, if the model identifies a segment of customers who have shown a drop in engagement with marketing content, the business could design specific outreach campaigns tailored to reignite interest and improve retention. Similarly, if the model indicates that certain purchasing behaviors are linked with higher churn rates, the business might consider offering personalized discounts or recommendations to keep those customers engaged.

2.2 FRAMEWORK OF LIFE CYCLE ASSESSMENT(LCA):

Life Cycle Assessment (LCA) is a structured framework used to assess the environmental impacts of a product, process, or service throughout its entire life cycle. As industries and organizations increasingly prioritize sustainability, LCA has become an essential tool for understanding and reducing environmental footprints. By evaluating every stage of a product's life—from raw material extraction to disposal—LCA enables companies and policymakers to make informed decisions about product design, material selection, manufacturing processes, and end-of-life options. This framework follows a systematic process organized into four primary stages: goal and scope definition, life cycle inventory, life cycle impact assessment, and interpretation.

1. Goal and Scope Definition

The first stage of LCA is to establish the goal and scope, which defines the purpose of the assessment, the functional unit, the system boundaries, and any assumptions or limitations. The goal articulates the primary objectives of the study—such as reducing emissions, comparing two products, or assessing compliance with environmental regulations. This stage also identifies the stakeholders involved and determines whether the LCA is intended for internal use or external reporting.

The functional unit is a critical aspect of this stage, as it establishes a standardized basis for comparison across products or services. For instance, if the assessment is for a beverage bottle, the functional unit might be "one liter of bottled water." System boundaries are defined to specify which life cycle stages are included in the assessment. Boundaries might include "cradle-to-grave" (covering every stage from raw material extraction to disposal), "cradle-to-gate" (from raw material extraction to the factory gate), or other scopes depending on the study's focus. This phase ensures that the LCA is aligned with the specific goals and remains relevant to the intended audience.

2. Life Cycle Inventory (LCI)

In the Life Cycle Inventory phase, data is collected on all the inputs and outputs associated with each stage of the product's life cycle. This includes raw material extraction, energy use, emissions, waste generation, and transportation at each step. The goal is to quantify the resources consumed (e.g., water, energy) and pollutants emitted (e.g., greenhouse gases, particulate matter) across the defined system boundaries. Inventory data is often extensive and requires careful consideration of data sources, quality, and completeness.

This stage is often the most resource-intensive part of an LCA as it involves gathering data from multiple sources, which can include direct measurements, databases, or estimations. Once the inventory is complete, the data provides a foundation for identifying the environmental hotspots—areas where resource consumption or emissions are particularly high—throughout the product's life cycle. These hotspots help direct attention to the most impactful stages of the life cycle, offering a basis for potential improvements.

3. Life Cycle Impact Assessment (LCIA)

The Life Cycle Impact Assessment phase evaluates the environmental impacts of the inputs and outputs quantified in the LCI. This involves translating raw data into meaningful impact categories such as global warming potential, ozone depletion, acidification, resource depletion, and toxicity. The LCIA phase helps link inventory data to specific environmental consequences, providing a clearer picture of the ecological footprint of each stage in the product's life cycle.

In LCIA, impact assessment models are applied to convert inventory data into potential environmental impacts. The method may involve characterization (quantifying the impact), normalization (comparing the impact to a reference), and weighting (prioritizing impact categories based on importance.

4. Interpretation

The final stage of the LCA framework is Interpretation, where the results from the previous stages are analyzed, summarized, and discussed in the context of the study's initial goals. This phase identifies significant environmental impacts, evaluates uncertainties, and assesses the reliability of the findings. Interpretation is an iterative process that involves reviewing and refining data, checking for consistency, and ensuring transparency.

Interpretation may include recommendations for improvement, highlighting which stages or processes should be optimized to minimize environmental impacts. For example, if the LCIA reveals that a product's manufacturing stage has a high carbon footprint, companies may explore alternative materials or more efficient production techniques to mitigate this impact. This stage also includes an evaluation of limitations and uncertainties within the assessment to provide context for the results. For instance, if data were based on estimates rather than direct measurements, this should be noted to inform the confidence level of the results.

CHAPTER 3

3. SYSTEM OVERVIEW

3.1 EXISTING SYSTEM

The existing system for predicting customer churn in e-commerce often relies on machine learning algorithms that include logistic regression, decision trees, and random forests.

Logistic Regression

Logistic regression is a popular choice due to its simplicity and efficiency, particularly for binary classification problems such as identifying whether a customer will churn or not. This algorithm models the probability of a binary outcome by examining relationships between a set of input features and the predicted output. While it is valued for its straightforward interpretation and relatively low computational cost, logistic regression is limited in its ability to capture complex, non-linear patterns in data. Moreover, it often struggles with imbalanced data, which can be common in churn prediction scenarios where the number of churned customers is significantly lower than non-churned ones.

Decision Trees:

Decision trees, on the other hand, offer a more intuitive approach. They segment data into branches based on feature values, forming a tree-like structure where each path represents a decision rule, leading to either a churn or non-churn outcome. The main advantage of decision trees lies in their visual simplicity and the ability to handle both numerical and categorical data seamlessly. This makes them suitable for capturing non-linear relationships between features and responses. However, decision trees have notable drawbacks, primarily their tendency to overfit. Overfitting occurs when the model becomes too complex and starts to memorize the training data, losing its ability to generalize well to new, unseen data. This instability means that even small changes in the input data can lead to a vastly different tree structure.

Random Forest:

Random forests enhance the predictive power by combining multiple decision trees into an ensemble. Each tree in the forest is trained on a random subset of the data, and their combined output is averaged to produce a more accurate and stable prediction. This method reduces the risk of overfitting seen in individual decision trees and generally improves the reliability of the model's predictions. Additionally, random forests can provide insights into which features are most influential, helping analysts understand what drives customer churn. However, the trade-off comes in the form of increased computational requirements. As the number of trees grows, so does the complexity and resource demand, making the system more challenging to interpret and slower to execute, especially when working with very large datasets.

Limitations of the Existing System

While these models provide solid foundations for churn prediction, they have certain drawbacks:

Despite the strengths of these existing models, there are persistent limitations that impact their effectiveness in real-world e-commerce applications. Logistic regression's reliance on linear relationships restricts its applicability to more complex data patterns. Decision trees, while straightforward, can produce models that are overly sensitive to data fluctuations, reducing their reliability. Random forests, although robust and capable of handling complex feature interactions, can be resource-intensive and require significant processing power for training and prediction, which can hinder their use in environments with limited computational resources.

Given these challenges, businesses have sought more advanced solutions to enhance prediction accuracy and model efficiency. The existing system's inability to handle categorical data without extensive preprocessing poses another significant challenge. In e-commerce, where categorical features such as product categories, customer segments, and regions play a crucial role, conventional methods can be cumbersome. This need for improved handling of categorical data, coupled with the goal of reducing overfitting and computational burden, has spurred the exploration of more sophisticated algorithms like CatBoost, which addresses these shortcomings by simplifying data preprocessing, managing missing values effectively, and providing high-accuracy predictions while maintaining computational efficiency.

These challenges point to the need for more advanced algorithms, such as **CatBoost**, which can address these issues by efficiently handling categorical data, managing missing values, and providing higher accuracy with reduced risk of overfitting.

3.2 PROPOSED SYSTEM

The proposed system aims to predict customer churn in e-commerce using a machine learning algorithm known as CatBoost. This system is designed to improve the accuracy of churn predictions, especially in managing categorical data, which is abundant in customer transactions and behavioral datasets. The model's primary goal is to aid e-commerce businesses in identifying at-risk customers and providing actionable insights that enable targeted retention strategies. Additionally, a comprehensive dashboard will offer visualizations and trends in customer churn, helping businesses monitor changes in customer loyalty and adapt their marketing strategies accordingly.

Key Features of the Proposed System

1. **CatBoost Algorithm Implementation**:

The proposed system employs the CatBoost algorithm, a gradient boosting algorithm specifically designed to handle categorical features and missing values, which are often present in e-commerce datasets. Unlike traditional algorithms like logistic regression or random forests, CatBoost is robust against overfitting and can process complex data relationships efficiently. This high accuracy and reliability make it a preferred choice for predicting churn, providing businesses with precise and actionable insights.

2. **Interactive Dashboard**:

The dashboard is a critical component of the system, enabling users to interact with the churn prediction data in real-time. Visualizations include key metrics such as churn probability, customer segments, and predictive insights, which are instrumental in understanding the reasons behind customer churn. Businesses can leverage this dashboard to segment their customer base, prioritize retention efforts, and measure the effectiveness of interventions.

3. **Report Generation**:

Users can generate downloadable reports summarizing churn predictions and insights. These reports offer valuable information for stakeholders to analyze customer behavior trends and evaluate long-term retention strategies. Additionally, the report feature streamlines the sharing of findings with teams, aiding in collaborative efforts to enhance customer loyalty.

Advantages of the Proposed System

The proposed system offers several significant advantages:

- **High Predictive Accuracy**: With CatBoost, the system delivers precise churn predictions, enhancing the reliability of insights.
- **Efficient Categorical Data Handling**: CatBoost simplifies data preprocessing by efficiently managing categorical features, which are prevalent in e-commerce datasets.
- **Reduction in Overfitting**: The algorithm's design reduces the risk of overfitting, ensuring the model generalizes well to new data, which is essential for maintaining prediction accuracy over time.

3.3 FEASBILITY STUDY

This feasibility study evaluates the potential for successfully implementing a machine learning-based customer churn prediction system for e-commerce. By leveraging the CatBoost algorithm, the proposed system aims to provide accurate churn predictions and insights that can help businesses improve customer retention. The system includes interactive dashboards and report generation features, ensuring stakeholders have access to actionable data on at-risk customers.

Economic Feasibility

The economic feasibility of the proposed customer churn prediction system is strong, as it promises to enhance customer retention rates, thereby boosting revenue and lowering the costs associated with customer acquisition. By identifying at-risk customers through precise churn predictions, businesses can allocate resources to targeted retention efforts, reducing the need for broad, costly marketing campaigns. Although initial development and implementation costs include expenses for data scientists, software engineers, and cloud infrastructure, these are offset by the system’s long-term impact on profitability. The return on investment (ROI) is expected to be high, with improved customer loyalty and retention leading to increased customer lifetime value (CLV) and a better allocation of marketing budgets.

Technical Feasibility:

The technical feasibility of this system is well-supported by its use of the CatBoost algorithm, which excels in handling categorical data and missing values commonly found in e-commerce datasets. The system requires scalable cloud infrastructure capable of managing large data volumes and complex computations, such as those offered by AWS or Google Cloud. Additionally, the system will use popular machine learning libraries and web frameworks to ensure compatibility, performance, and ease of integration with existing e-commerce platforms. Risks like model complexity and data integration challenges can be managed with thorough data preprocessing and robust data management practices, ensuring the model's accuracy and responsiveness as data scales.

Operational Feasibility

Operationally, the system is designed for smooth integration into e-commerce business processes, with a user-friendly dashboard that allows marketing and customer service teams to interpret predictions easily. The dashboard's intuitive visualizations, combined with automated report generation, ensure that stakeholders can quickly act on churn insights without extensive technical training. Minimal adjustments in organizational workflows are needed, as the system aligns with standard business goals around customer engagement and retention. Training sessions for non-technical staff will further support adoption, enabling teams to leverage the system's predictions in daily retention activities, making it a practical addition to operational processes.

In summary, the feasibility study indicates that implementing the e-commerce customer churn prediction system is both viable and beneficial for businesses. From an economic perspective, the system promises a high ROI through improved customer retention and targeted marketing. Technically, the choice of the CatBoost algorithm, combined with scalable cloud infrastructure, ensures that the system can handle large data volumes effectively. Operationally, a user-friendly dashboard and comprehensive reporting facilitate integration into business processes, while organizational support can drive smooth change management.

CHAPTER 4

4. SYSTEM REQUIREMENTS

4.1 HARDWARE REQUIREMENTS

1. Core System Requirements

Processor: Intel Core i7 (10th Gen or higher) or AMD Ryzen 7 (or equivalent)

- Multi-core processors enhance parallel processing, which is essential for large-scale data analysis and model training.

Memory (RAM): 16 GB or more

- Sufficient memory is necessary to load and process large datasets without slowdowns.

2. Additional Components

Operating System: Windows 10/11, macOS, or a Linux distribution (e.g., Ubuntu 20.04)

- Compatible with Python and machine learning frameworks.

3. Networking and Backup

Internet Connection: 10 Mbps or higher

- Reliable internet for cloud storage, collaboration, and remote model deployment.

Backup Storage: External Hard Drive (1 TB) or Cloud Storage (e.g., Google Drive, AWS S3)

- Necessary for regularly backing up datasets and trained models.

4.2 SOFTWARE REQUIREMENTS

1. Frontend Development

- **ReactJS**: A JavaScript library for building user interfaces, especially single-page applications.
    - **Version**: 17.0 or later
    - **Features**:
        - Component-based architecture, facilitating code reuse.
        - Virtual DOM for efficient UI updates.
        - Support for JSX, making it easier to write HTML in JavaScript.
- **React Router**: For handling routing in the single-page application.
    - **Version**: 6.0 or later
    - **Purpose**: Enables navigation between different views within the application without reloading the page.
- **Axios**: A promise-based HTTP client for making API requests from the frontend to the backend.
    - **Version**: Latest stable version
    - **Purpose**: Facilitates communication with the NodeJS backend, especially for CRUD operations related to customer data and churn predictions.

2. Backend Development

- **NodeJS**: JavaScript runtime environment for building the server-side of the application.
    - **Version**: 14.0 or later
    - **Features**:
        - Non-blocking I/O operations for handling multiple requests simultaneously.
        - Compatibility with JavaScript on both frontend and backend, simplifying development.

- **ExpressJS**: Minimalist framework for building APIs with NodeJS.
    - **Version**: 4.17 or later
    - **Purpose**: Simplifies API creation and request handling, allowing efficient management of routes, middleware, and controllers.
- **CatBoost Model Integration**:
    - **Model File**: Serialized CatBoost model file (e.g., in .cbm or .pkl format).
    - **Purpose**: The backend will load this pre-trained model to predict churn for new data inputs.
    - **Integration**: Load the model using Python in a separate microservice or integrate with REST API to interact with the NodeJS backend.

3. Database

- **MongoDB**: NoSQL database for storing customer data, churn predictions, and engagement metrics.
    - **Version**: 4.0 or later
    - **Features**:
        - Schema-less data storage, allowing flexibility for storing varied customer data.
        - JSON-like document structure compatible with JavaScript-based NodeJS and ReactJS.
        - High scalability and performance for handling large datasets.
- **Mongoose**: ODM (Object Data Modeling) library for MongoDB and NodeJS.
    - **Purpose**: Provides a straightforward schema-based solution for interacting with MongoDB, enabling data validation and querying.

4. Machine Learning and Data Processing

- **Python (3.7 or later)**: Used for training the CatBoost model and data preprocessing.

    - **Packages**:
        - **CatBoost**: For model training and prediction.
        - **scikit-learn**: For additional preprocessing and evaluation tools.
        - **Pandas, NumPy**: For data manipulation and analysis.
- **REST API or Flask API** (optional): Interface to serve the Python-based machine learning model, allowing the NodeJS backend to make API calls to retrieve churn predictions.
    - **Purpose**: Separates the machine learning workload and enables scalability by isolating the model in a microservice.

5. Development Tools and Dependencies

- **NPM (Node Package Manager)**: For managing JavaScript libraries and dependencies.
    - **Purpose**: Facilitates installing libraries (e.g., ExpressJS, Mongoose, Axios) and running scripts for building, testing, and deploying the application.
- **Git**: Version control system for managing source code.
    - **Features**: Enables collaboration, version tracking, and code management. Compatible with GitHub, GitLab, and Bitbucket for code hosting.

6. Testing and Deployment

- **Postman**: API testing tool.
    - **Purpose**: Used for testing REST API endpoints on the backend, including model prediction endpoints and database interactions.
- **Docker** (optional): Containerization tool.
    - **Purpose**: Enables consistent deployment across various environments by packaging the application and its dependencies into containers.
- **Cloud Hosting** (e.g., AWS, Google Cloud, Heroku):

- **Purpose**: Deploy the application on a cloud provider to make it accessible online, with scalability options to handle increased data volume.

4.3 FUNCTIONAL AND NON-FUNCTIONAL REQUIREMENTS

The e-commerce customer churn prediction system requires a robust set of functional and non-functional requirements to ensure it meets business needs, performs efficiently, and remains user-friendly. Below is a detailed outline of these requirements.

Functional Requirements

1. **User Authentication and Authorization**:
   - The system must support secure login mechanisms, including multi-factor authentication, to ensure only authorized users access sensitive customer data.
   - Role-based access control should be implemented, allowing different user roles (e.g., administrator, analyst, and manager) to have appropriate access privileges based on their responsibilities. This feature protects data privacy while ensuring relevant personnel can access necessary functionalities.
2. **Data Upload**:
   - Users should be able to upload customer datasets in commonly used formats, such as CSV and Excel. The upload feature must include validation checks to verify data quality, format consistency, and the presence of required fields (e.g., customer ID, purchase history, engagement metrics).
   - The system should handle large datasets, with capacities for up to 100,000 records per upload, ensuring compatibility with the data volumes common in e-commerce businesses.
3. **Data Preprocessing**:
   - The system should automatically preprocess uploaded data, handling missing values by imputing or removing them, and encoding categorical variables for model compatibility. This

preprocessing step is essential for ensuring accurate predictions and simplifying data preparation for users.

- Users should be notified of any data inconsistencies or errors that may impact model performance, with options to rectify them before running predictions.

4. **Churn Prediction**:
   - The core functionality of the system is to predict customer churn using a machine learning model, specifically the CatBoost algorithm. Once the data is processed, the model should generate churn probabilities for each customer, indicating their likelihood of disengaging from the platform.
   - These predictions should be generated with high accuracy and displayed in a way that allows users to identify at-risk customers and focus retention efforts on these segments.
5. **Dashboard and Visualization**:
   - The system must include a dashboard that visually displays churn metrics, customer segments, and trends over time. Visualizations such as bar charts, line graphs, and heatmaps should present churn data in an intuitive and easy-to-interpret format.
   - The dashboard should allow users to filter and segment customers based on factors such as churn probability, customer value, and engagement level, providing tailored insights that support targeted marketing strategies.
6. **Report Generation**:
   - Users should be able to generate detailed churn reports that summarize prediction results, customer segments, and insights into churn trends. These reports should be exportable in PDF and Excel formats for ease of sharing and offline analysis.
   - The report generation feature should include customization options, allowing users to select specific timeframes, metrics, or customer segments to focus on, enhancing the relevance of insights for different stakeholders.

Non-Functional Requirements

1. **Performance**:
    - The system should process datasets of up to 100,000 records within a maximum of two minutes to ensure efficiency. This performance standard allows users to upload large datasets and receive quick results, enhancing user experience and minimizing delays in decision-making.
    - Prediction latency should also be low, enabling real-time interaction with the dashboard and a responsive user experience when filtering or segmenting data.
2. **Scalability**:
    - The system should be scalable to handle increased data volumes and user demand as the business grows. The system architecture should support seamless scaling of both data processing and storage capacities to accommodate a larger user base and growing datasets without performance degradation.
    - Cloud-based infrastructure options, like AWS or Google Cloud, are recommended for their flexibility in scaling resources up or down as needed, allowing the system to grow cost-effectively.
3. **Security**:
    - Security is critical, as the system handles sensitive customer data. The system should use encryption protocols to protect data in transit and at rest, ensuring customer information remains confidential and protected from unauthorized access.
4. **Usability**:
    - The system should be intuitive and user-friendly, with a clear interface and visual feedback for each action. Visual aids, like charts and graphs, simplify data interpretation, while tooltips, guides, and error messages should support users in navigating the system easily.
    - A minimal learning curve should enable non-technical users, such as marketing and customer support teams, to interact with the dashboard and interpret churn predictions without extensive training.

5. **Availability and Reliability**:
    - The system should ensure 99.9% uptime to support continuous access for users, especially during peak hours when data uploads or churn insights may be critical for decision-making.
    - Backup and disaster recovery mechanisms should be implemented to maintain system reliability and prevent data loss, ensuring that the system remains available during unexpected outages or server issues.
6. **Data Integrity and Accuracy**:
    - Ensuring data accuracy throughout the system is essential, as inaccurate or inconsistent data can impact prediction quality. The system should implement validation checks during data upload and preprocessing to ensure data integrity, identifying and handling anomalies automatically.
    - Version control mechanisms can be included to track model updates and maintain historical prediction data, providing users with confidence in the model's reliability and consistency.
7. **Maintainability and Extensibility**:
    - The system should be modular and maintainable, allowing for easy updates, bug fixes, and feature additions. A well-documented codebase with clean architecture can support future improvements and the addition of new features, such as integrating other machine learning algorithms or additional data sources.
    - Extensibility also includes the capability to incorporate new datasets or adapt to changes in customer behavior, ensuring the system remains relevant and useful in dynamic business environments.

CHAPTER 5

SYSTEM DESIGN

5.1 SYSTEM ARCHITECTURE:

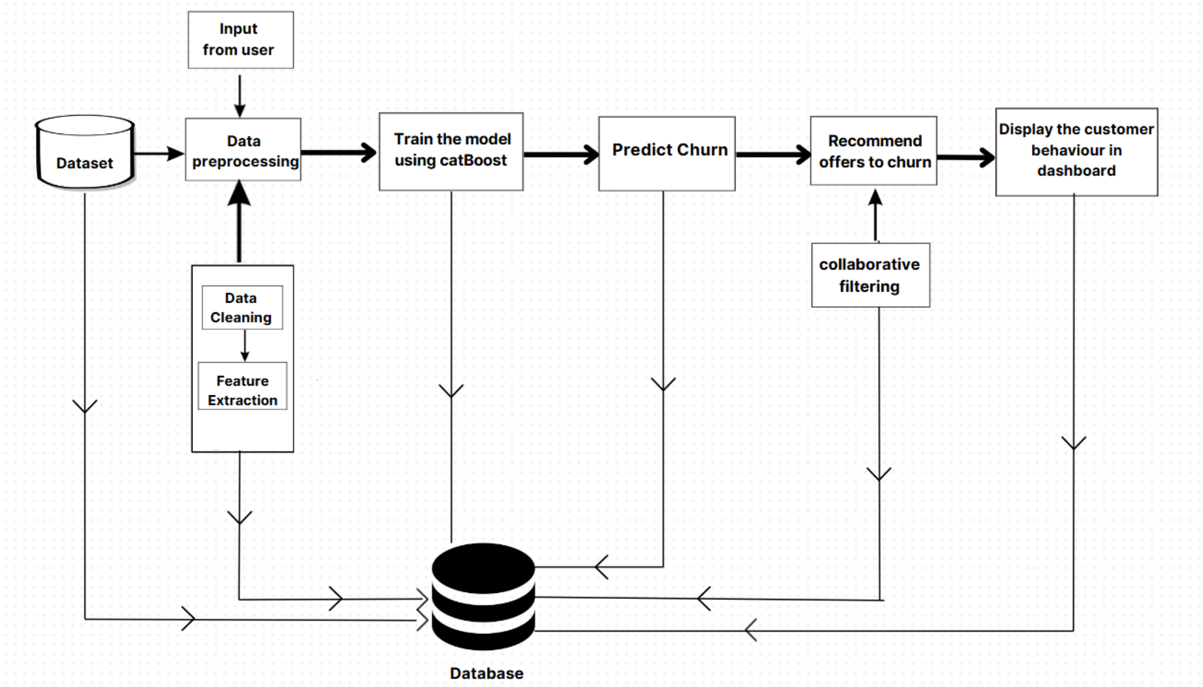

Fig 1-Architecture Diagram

Overview

The goal of this architecture is to predict customer churn in an e-commerce setting and generate personalized recommendations to reduce the likelihood of customer attrition. This system is designed in a modular and data-driven fashion to ensure accuracy, scalability, and actionable insights. The primary components of this architecture are data preprocessing, model training, churn prediction, recommendations, and visualization through a dashboard. Here's a step-by-step overview of each component.

1. Data Collection and Data Storage

The system starts by gathering relevant data on customer behavior, including demographics, purchase history, engagement metrics, and transaction data. All data collected is stored in a centralized **Database**. This database is critical to the entire architecture, as it supports the data preprocessing, training, and recommendation steps by providing necessary information for model training and decision-making.

- **Dataset**: This is the source of raw data, which includes both historical customer data and real-time updates. It contains fields such as customer demographics, transaction history, engagement data, and interaction logs.
- **Database**: Acts as the primary storage for all processed data and intermediate outputs. It also stores the results of predictions and recommendations, allowing for easy retrieval when displaying insights on the dashboard.

2. Data Preprocessing

Data preprocessing is a crucial step that involves preparing the raw data for analysis. This phase includes two main tasks: **Data Cleaning** and **Feature Extraction**.

- **Data Cleaning**: Involves removing inconsistencies, handling missing values, and correcting data anomalies. Clean data ensures that the model can be trained effectively without biases caused by errors or noise in the data.
- **Feature Extraction**: Here, relevant features are derived from the data to improve model accuracy. Examples include calculating RFM (Recency, Frequency, and Monetary value) scores, customer lifetime value, and engagement metrics. These features are crucial in determining the likelihood of churn.

3. Model Training Using CatBoost

In this architecture, the **CatBoost** algorithm is selected for training the churn prediction model. CatBoost is a powerful gradient-boosting algorithm that is particularly effective at handling categorical data, making it well-suited for customer churn prediction.

- **Training Process**: The model is trained on historical data, using features generated in the preprocessing stage. During training, the model learns patterns that correlate with churn by analyzing past customer behavior and churn outcomes.
- **Model Validation**: To ensure the model's robustness, it is validated using a test dataset, where metrics such as accuracy, precision, recall, and F1-score are calculated to evaluate its predictive power.

The trained model is then saved and used to make predictions on current customer data. This model is continuously updated with new data to improve its predictive accuracy over time.

4. Churn Prediction

Once the model is trained, it is deployed to predict the churn likelihood for each customer.

- **Input from User**: Current customer data is fed into the model, which generates a churn prediction. This input could be real-time customer data or periodic batch updates, depending on the system's operational needs.
- **Predict Churn**: The model evaluates each customer based on their attributes and interactions, assigning a churn probability score. Customers are then segmented based on their churn risk (e.g., low, medium, or high risk). This segmentation allows for targeted intervention, focusing resources on high-risk customers.

5. Recommendation Engine Using Collaborative Filtering

Once churn predictions are made, the system initiates the **Recommendation Engine** to provide actionable insights on retaining at-risk customers. Collaborative filtering is used to generate personalized recommendations that encourage engagement and reduce churn risk.

- **Collaborative Filtering**: This recommendation technique analyzes data from similar users who did not churn to suggest relevant offers, products, or services to at-risk customers. By leveraging historical interactions and preferences, the system tailors recommendations for each customer segment based on what has worked with similar users in the past.
- **Personalized Offers**: For high-risk customers, personalized discounts, loyalty rewards, or exclusive offers can be generated to encourage continued engagement. The recommendation engine identifies the best strategies for each segment, ensuring that interventions are both timely and effective.

The recommendations generated by the system are then saved in the database, allowing for easy access by the dashboard.

6. Dashboard for Visualization and Insights

The final component of the system is the **Dashboard**, which serves as the user interface for visualizing customer behavior, churn predictions, and recommendations.

- **Display Customer Behavior**: The dashboard provides detailed insights into customer behavior and segmentation based on churn risk scores. This includes visualizations such as churn risk distribution, customer activity trends, and engagement metrics.

- **Recommendation Display**: Recommendations generated by the collaborative filtering engine are displayed, allowing customer service and marketing teams to see the suggested actions for each at-risk customer.
- **Action Tracking**: The dashboard can also track the effectiveness of recommendations, allowing users to monitor which actions are most successful in reducing churn. This feedback loop helps refine future recommendations, improving the system's accuracy and relevance over time.

5.2 MODULES DESCRIPTION

5.2.1 DATA PREPROCESSING MODULE

Objective: To prepare the raw customer data for modeling by cleaning, transforming, and encoding it into a format suitable for analysis.

Key Steps:

- **Data Cleaning:**
  - Handle missing values: Use imputation techniques (e.g., mean, median, or mode imputation for numerical features, or mode for categorical features).
  - Remove or correct inconsistent data entries (e.g., invalid or duplicate rows).
  - Identify and handle outliers (e.g., via z-score or IQR methods).
- **Feature Engineering:**
  - Create new features from existing data (e.g., aggregate features like average purchase frequency, total spend, etc.).
  - Extract relevant features from the purchase history (e.g., recency, frequency, and monetary value—RFM analysis).
  - Convert timestamps to meaningful features such as "days since last purchase" or "customer lifetime."
- **Categorical Encoding:**
  - Use Label Encoding or One-Hot Encoding to convert categorical variables (e.g., product category, payment method) into numerical format for model training.

    - CatBoost can handle categorical features directly, so they should be labeled appropriately for the model.
- **Feature Scaling:**
    - For algorithms like Logistic Regression or SVM, ensure features are standardized or normalized, but CatBoost generally handles unscaled data well.

Deliverables:

- A cleaned and transformed dataset ready for modeling.
- A report detailing the preprocessing steps, including data cleaning techniques, feature engineering, and encoding strategies.

5.2.2 PREDICTION MODULE

Objective: To build and train the CatBoost model to predict customer churn and evaluate its performance.

Key Steps:

- **Model Selection:**
    - Use CatBoost (a gradient boosting algorithm) for churn prediction due to its ability to handle categorical variables and missing values effectively.
    - Specify parameters for CatBoost like learning rate, number of trees, depth of trees, etc.
- **Data Splitting:**
    - Split the dataset into training and testing sets (e.g., 80% training and 20% testing).
    - Optionally, use cross-validation to ensure robust performance evaluation.
- **Model Training:**
    - Train the CatBoost model on the training data using the appropriate features (e.g., demographics, purchase history, engagement metrics).
    - Tune hyperparameters for optimal performance using Grid Search or Random Search.
- **Prediction:**
    - Use the trained model to predict churn for customers in the testing set.
    - Output predicted probabilities of churn (e.g., likelihood of a customer churning within a given timeframe).

5.2.3 EVALUATION MODULE

Objective: To evaluate the performance of the churn prediction model and assess its effectiveness in predicting customer churn.

Key Steps:

- **Performance Metrics:**
  - Accuracy: Percentage of correctly predicted churn/non-churn cases.
  - Precision: The proportion of true positives out of all predicted positives.
  - Recall: The proportion of true positives out of all actual positives (important for churn prediction as we want to minimize false negatives).
  - F1-Score: The harmonic mean of precision and recall.
  - AUC-ROC: Area under the receiver operating characteristic curve to evaluate classification performance across all thresholds.
  - Confusion Matrix: Visualize the model's true positives, false positives, true negatives, and false negatives.
- **Feature Importance:**
  - Extract and display the most important features in the model, helping to understand which customer behaviors and metrics contribute most to churn prediction.
- **Model Comparison:**
  - Optionally compare the CatBoost model against other models (e.g., Logistic Regression, Random Forest) in terms of performance metrics and feature importance.

Deliverables:

- **Evaluation metrics** (e.g., accuracy, precision, recall, F1-score).

5.2.4 WEB INTERFACE MODULE

Objective: To build a user-friendly web interface or dashboard that visualizes churn prediction results and helps stakeholders interact with the model and insights.

Key Features:

- **Dashboard Overview:**
  - Display high-level metrics such as total churned customers, churn rate, and predicted churn probabilities for each customer.

  - Visual representations of the churn predictions (e.g., bar charts, pie charts, or line graphs showing churn trends over time).
- **Customer Segmentation:**
  - Provide segmentation options for businesses to view churn predictions by customer demographics, behavior, or purchase history.
  - Allow filtering by specific timeframes (e.g., "churned customers in the last 30 days").
- **Retention Insights:**
  - Visualize factors contributing to churn (e.g., low engagement, high return rates, infrequent purchases).
  - Suggest retention strategies based on the model's findings (e.g., targeted promotions, personalized offers, engagement strategies).
- **Report Download:**
  - Provide a feature for stakeholders to download detailed reports, including churn predictions, customer segments, and analysis.
- **Interactive Model Input:**
  - Allow users to input customer details (e.g., age, frequency of purchases, product categories) to predict churn for individual customers.

6 UML DIAGRAM

6.1 DATA FLOW DIAGRAM:

LEVEL 0:

The DFD depicts a system that processes user-uploaded data to predict customer churn. The user logs in, uploads data, which is preprocessed, and then used to predict churn. The results are visualized on a dashboard. The system likely involves a user database and stores both uploaded and processed data. This system could be valuable for businesses to proactively retain customers.

LEVEL1:

The data flow for the churn prediction system begins at the "Start" node, representing the initialization of the process. The flow then moves to the "User" block, where end-users provide input data such as customer information, usage patterns, and other relevant metrics. This data is sent to the "Churn Prediction System," the core component that processes the input using machine learning algorithms to predict customer churn. The system analyzes patterns and calculates the likelihood of customers leaving based on various features. The results are then passed to the "Database" block, where all relevant data—including input information and prediction results—are stored for future use, such as reporting or model retraining. Finally, the process concludes at the "End" node, signifying the completion of the churn prediction workflow. This left-to-right flow ensures that each component processes data sequentially, from start to end.

LEVEL2:

The DFD depicts a system for predicting customer churn. Users log in, upload data, which is validated, cleaned, and transformed. Machine learning models are used to generate churn predictions, visualized through dashboards and reports. The system manages user data, stores predictions, and provides insights for customer retention strategies.

6.2 USECASE DIAGRAM:

Actors: Customer, System

1. Upload Dataset: The customer uploads a dataset containing relevant information for churn prediction.

2. Run Churn Prediction: The system processes the uploaded dataset to predict churn probabilities.
3. View Prediction Results: The customer views the results of the churn prediction, which displays insights on likely customer churn.
4. Data Analysis with Dashboard: The customer can analyze the churn data using a dashboard that provides various analytical insights and visualizations.
5. Download Report: The customer has the option to download a report containing detailed churn analysis results for offline reference.

6.3 CLASS DIAGRAM:

This UML class diagram represents the structure and relationships between different components in a churn prediction application.

1. Dataset: Handles the datasets, including attributes like datasetID, fileName, and filePath, and methods for uploading datasets, performing analysis, viewing results, and downloading reports.
2. User: Represents users with attributes like userID, username, and email. It includes methods for uploading datasets, viewing results, performing analysis, and downloading reports.
3. ChurnPredictionModel: Manages the churn prediction model with attributes such as modelID, algorithm, and accuracy. Methods include running predictions on datasets and generating reports.
4. PredictionResult: Stores prediction results with attributes such as resultID, churnRate, predictionDetails, and timestamp. It includes methods for displaying results and exporting reports.
5. Dashboard: Represents the user's dashboard, which includes attributes for dashboardID, userID, and widgets. It has methods to display visualizations, perform analysis on datasets, and update results based on prediction output.

CHAPTER 7

SOFTWARE DEVELOPMENT LIFE CYCLE MODEL

AGILE MODEL

For this project, we implemented an Agile methodology to develop a customer churn prediction application, allowing for flexibility, iterative improvements, and collaboration across teams. The project was structured around sprints, each lasting two weeks, where deliverables included core functionalities like user authentication, data upload, machine learning predictions, and data visualization. In initial sprints, backend API structures were created to handle authentication, data upload, and prediction requests, while a trained CatBoost model API was deployed in Python to provide churn predictions based on customer data. Simultaneously, the frontend team developed a React-based interface to streamline the user experience, including components for uploading data, viewing predictions, and accessing reports. Subsequent sprints involved integrating the frontend with backend services using REST APIs, adding error handling, and enhancing user experience through a dashboard that visualizes churn insights using `chart.js`. Regular sprint reviews enabled us to gather stakeholder feedback, prioritize new features, and refine requirements, ensuring alignment with business goals. Testing was integrated continuously, with JUnit for backend endpoints and unit tests for key frontend components. Finally, we conducted user acceptance testing (UAT) to validate functionality and ensure robustness before the application launch. This Agile approach helped in accommodating feedback, improving cross-functional collaboration, and delivering a dynamic, user-centric solution effectively.

CHAPTER 8

PROGRAM CODE AND OUTPUTS

8.1 SAMPLE CODE

```
import React, { useState } from 'react';
import axios from 'axios';
import './UploadData.css';

function UploadData() {
  const [file, setFile] = useState(null);
  const [status, setStatus] = useState("");
  const [error, setError] = useState("");

  const handleFileChange = (e) => {
    setFile(e.target.files[0]);
    setStatus("");
    setError("");
  };

  const handleUpload = async () => {
    if (!file) return alert("Please upload a CSV file.");
    const formData = new FormData();
    formData.append("file", file);

    try {
      setStatus("Uploading...");
      await axios.post('http://localhost:5000/api/upload', formData);
      setStatus("File uploaded successfully!");
    } catch (err) {
      console.error("Upload error", err);
      setError("Failed to upload file. Please try again.");
    }
  };

  return (
    <div className="upload-container">
      <h1>Upload Customer Data</h1>
      <div className="input-container">
        <label>Upload CSV File</label>
        <input type="file" accept=".csv" onChange={handleFileChange} />
      </div>
      <button onClick={handleUpload}>Upload</button>

      {status && <p className="status">{status}</p>}
      {error && <p className="error">{error}</p>}
    </div>
  );
}

export default UploadData;

import React, { useEffect, useState } from 'react';
import { Bar, Pie } from 'react-chartjs-2';
import axios from 'axios';
import './Dashboard.css';
import {
  Chart as ChartJS,
  CategoryScale,
  LinearScale,
  BarElement,
  Title,
  Tooltip,
  Legend,
  ArcElement
} from 'chart.js';

// Register necessary Chart.js components
ChartJS.register(CategoryScale, LinearScale, BarElement, Title, Tooltip, Legend, ArcElement);

function Dashboard() {
  const [churnData, setChurnData] = useState([]);
  const [error, setError] = useState("");

  useEffect(() => {
    const fetchChurnData = async () => {
      try {
         const response = await axios.get('http://localhost:5000/api/churnData');  // Fetch churn-related data
        setChurnData(response.data);
      } catch (err) {
        console.error("Error fetching churn data", err);
        setError("Failed to load churn data.");
      }
    };

    fetchChurnData();
  }, []);

  // Example data processing for visualization
  const churnCounts = churnData.reduce(
    (counts, item) => {
      if (item.prediction === 1) counts.churn += 1;
      else counts.noChurn += 1;
      return counts;
    },
    { churn: 0, noChurn: 0 }
  );

  // Bar chart data
  const barData = {
    labels: ['Churn', 'No Churn'],
    datasets: [
      {
        label: 'Customer Churn Prediction',
        data: [churnCounts.churn, churnCounts.noChurn],
        backgroundColor: ['#ff6384', '#36a2eb'],
      },
    ],
  };
```

```
// Pie chart data
const pieData = {
  labels: ['Churn', 'No Churn'],
  datasets: [
    {
      label: 'Churn Distribution',
      data: [churnCounts.churn, churnCounts.noChurn],
      backgroundColor: ['#ff6384', '#36a2eb'],
    },
  ],
};

return (
  <div className="dashboard-container">
    <h1>Dashboard</h1>
    {error && <p className="error">{error}</p>}

    {churnData.length > 0 ? (
      <div className="charts">
        <div className="chart">
          <h2>Churn Prediction Distribution</h2>
          <Bar data={barData} options={{ responsive: true }} />
        </div>
        <div className="chart">
          <h2>Churn Prediction Pie Chart</h2>
          <Pie data={pieData} options={{ responsive: true }} />
          </div>
        </div>
      ) : (
        <p>Loading data...</p>
      )}
    </div>
  );
}
export default Dashboard;

// src/App.test.js
import React from 'react';
import { render, screen } from '@testing-library/react';
import '@testing-library/jest-dom';
import App from './App';

test('renders welcome message', () => {
  render(<App />);

  // Check for the welcome message
  const welcomeElement = screen.getByText(/Welcome to the Customer Churn Prediction Platform/i);
  expect(welcomeElement).toBeInTheDocument();
});

test('renders Go to Churn Prediction button', () => {
  render(<App />);
  // Check for the presence of the "Go to Churn Prediction" button
  const buttonElement = screen.getByRole('button', { name: /Go to Churn Prediction/i });
  expect(buttonElement).toBeInTheDocument();
});
```

8.2 OUTPUT SCREENSHOTS

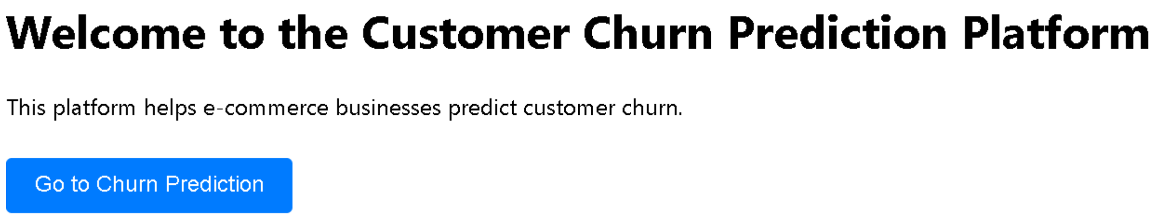

Welcome to the Customer Churn Prediction Platform

This platform helps e-commerce businesses predict customer churn.

Fig 2: web interface

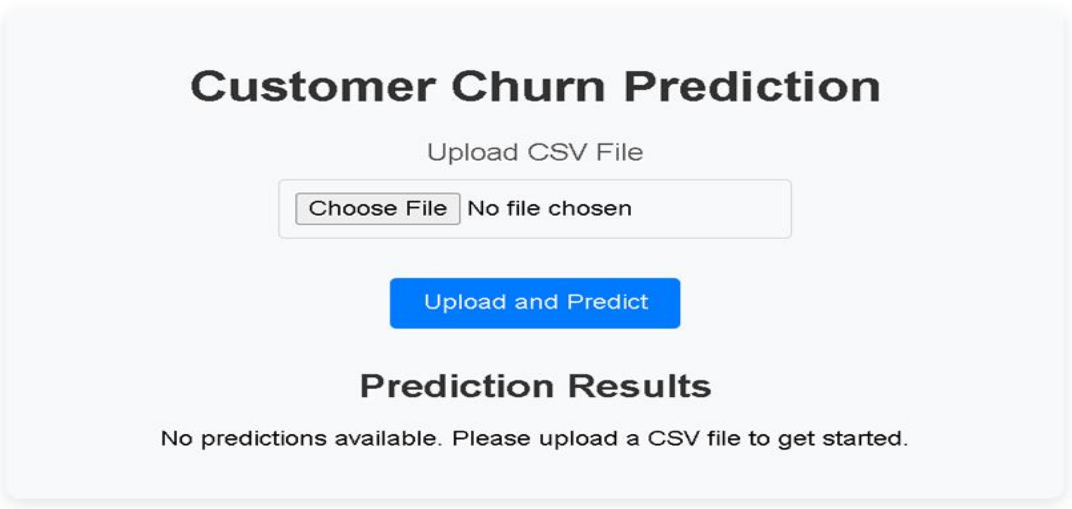

Customer Churn Prediction

Upload CSV File

Prediction Results

No predictions available. Please upload a CSV file to get started.

Fig 3: Prediction interface

DATA FLOW DIAGRAM:

LEVEL0:

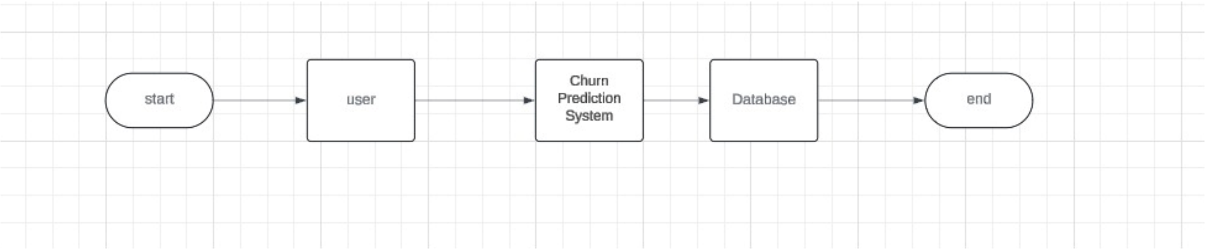

Fig 4: DFD LEVEL 0

LEVEL1:

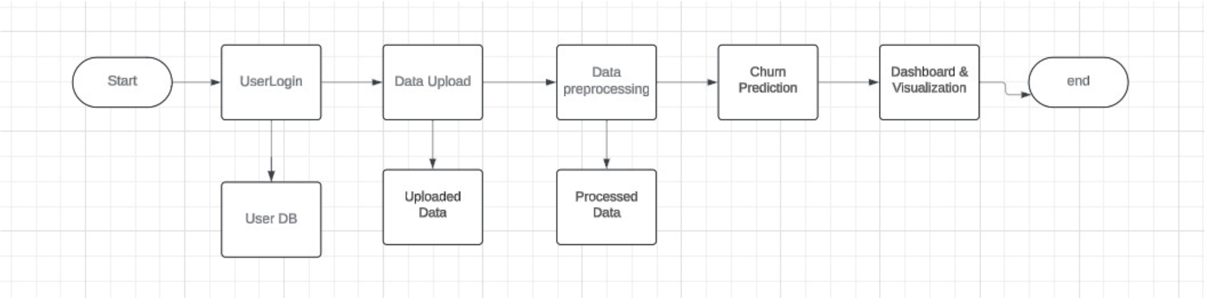

Fig 5: DFD LEVEL 1

LEVEL2:

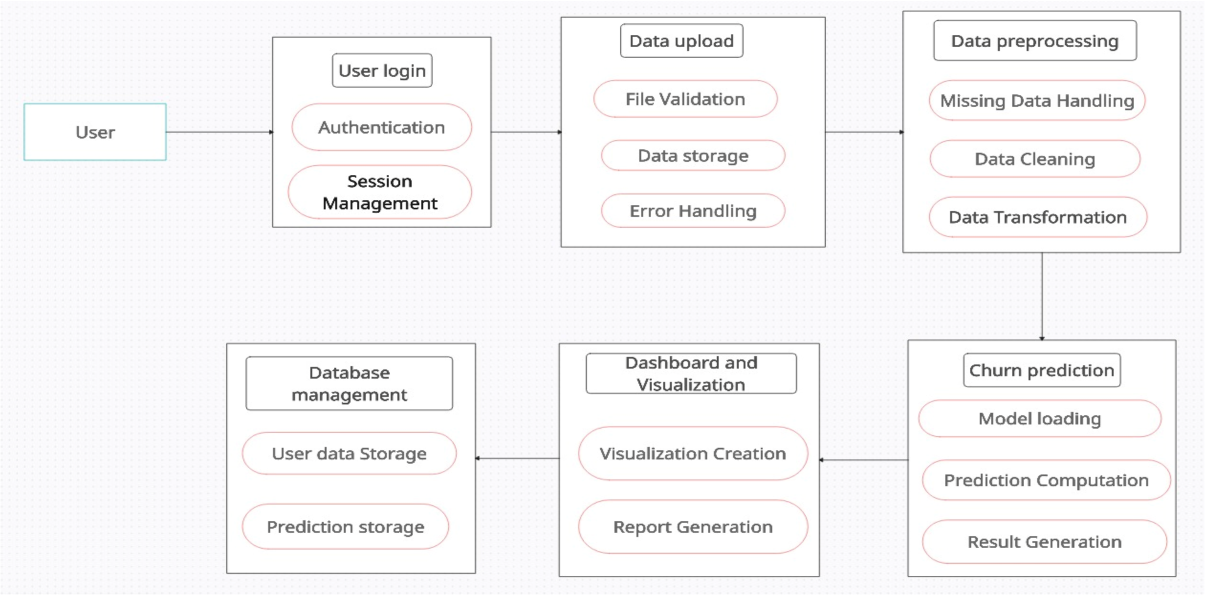

Fig 6: DFD LEVEL 2

USECASE DIAGRAM:

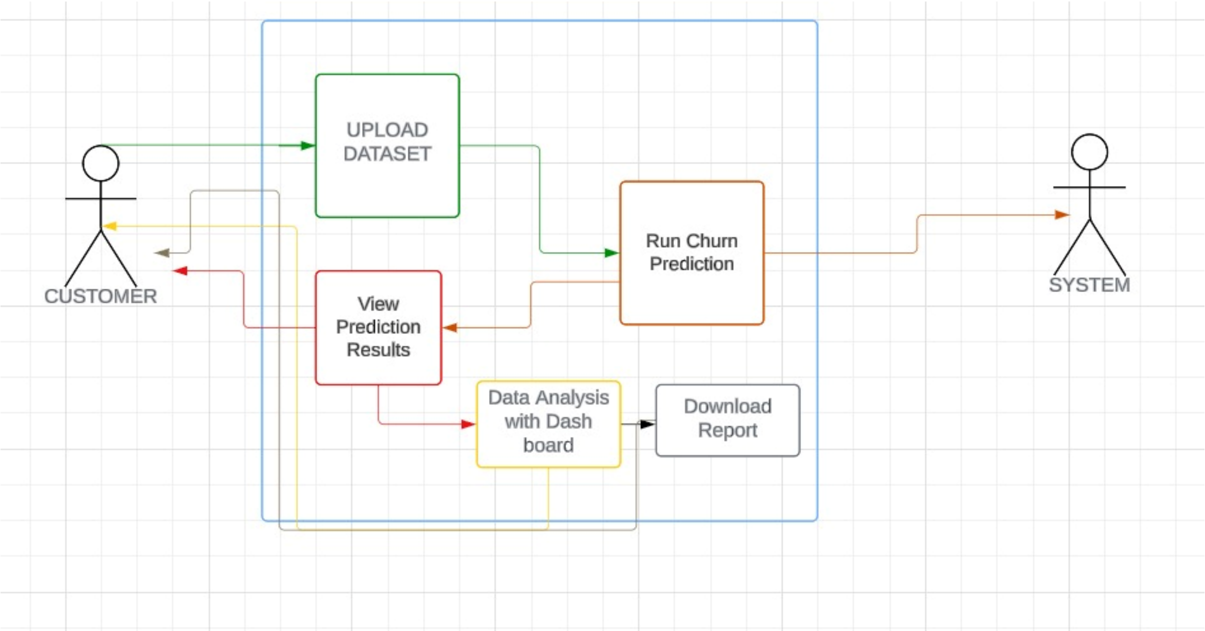

Fig 7:Use Case Diagram

CLASS DIAGRAM:

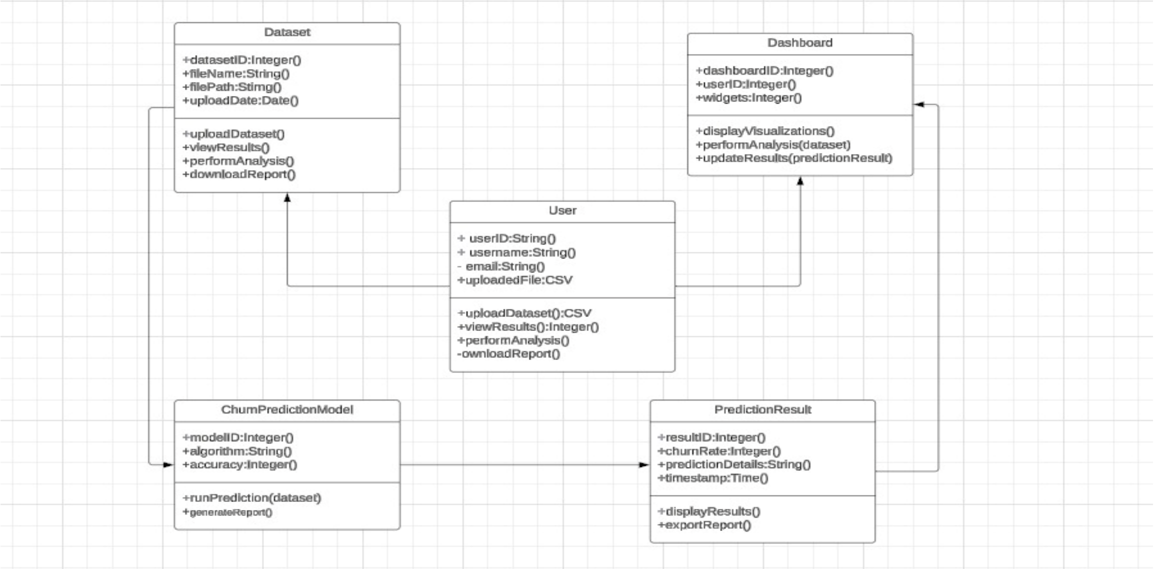

Fig 8: Class Diagram

CHAPTER 9

TESTING

9.1 FUNCTIONAL TESTING

Purpose: Verify that each function of the system performs according to the specified requirements.

Examples of Functional Tests:

Data Input Validation: Ensure the system correctly handles and validates various types of input data, including purchase history, browsing patterns, and engagement metrics.

Churn Prediction: Verify that the churn prediction model (CatBoost) accurately processes input data and provides a churn probability output.

Dashboard Functionality: Check that the dashboard accurately displays churn metrics, customer segments, and other predictive insights as designed.

User Interface (UI): Test that users can easily navigate and a

ccess different sections of the dashboard, filter data, and view detailed insights.

Alerts and Notifications: If there are automated alerts for high churn-risk customers, verify that these are correctly triggered based on prediction thresholds.

Outcome: The system should meet all functional requirements, ensuring that each component works correctly and provides accurate outputs.

9.2 UNIT TESTING:

Purpose: Validate individual components of the system in isolation to ensure each unit functions correctly.

Examples of Unit Tests:

Data Preprocessing Functions: Test each function responsible for data cleaning, transformation, and handling missing values.

Feature Engineering: Ensure that specific features (e.g., engagement metrics or purchase frequency) are correctly calculated and formatted before feeding into the model.

Churn Prediction Model: Verify that the CatBoost model functions as expected, producing predictions when given sample input data.

Dashboard Components: Test individual components of the dashboard (e.g., graphs, tables) to ensure they render correctly with sample data.

Utility Functions: For any helper functions (e.g., date conversions, data aggregations), verify that they provide the expected output for given inputs.

Outcome: Each unit of code should be validated for correctness, ensuring minimal bugs in individual components.

9.3 PERFORMANCE TESTING

Purpose: Assess the system's performance under various conditions to ensure it can handle real-world usage.

Examples of Performance Tests:

Load Testing: Simulate multiple users accessing the dashboard simultaneously to test its responsiveness and ensure it can handle high traffic without lag.

Model Prediction Time: Measure the time taken by the CatBoost model to process data and generate churn predictions, ensuring predictions are provided within an acceptable timeframe.

Data Processing Speed: Test how quickly the system processes large datasets to ensure timely and efficient data transformation and feature extraction.

Dashboard Rendering: Ensure that the dashboard loads quickly, with all charts and metrics rendering within a reasonable time, even with large volumes of data.

Scalability: Evaluate the system's performance when scaling the data size, to check if it can handle the growth of customer data over time.

Outcome: The system should perform optimally under expected and peak loads, ensuring it is efficient and responsive during usage.

9.4 TESTING SCREENSHOTS

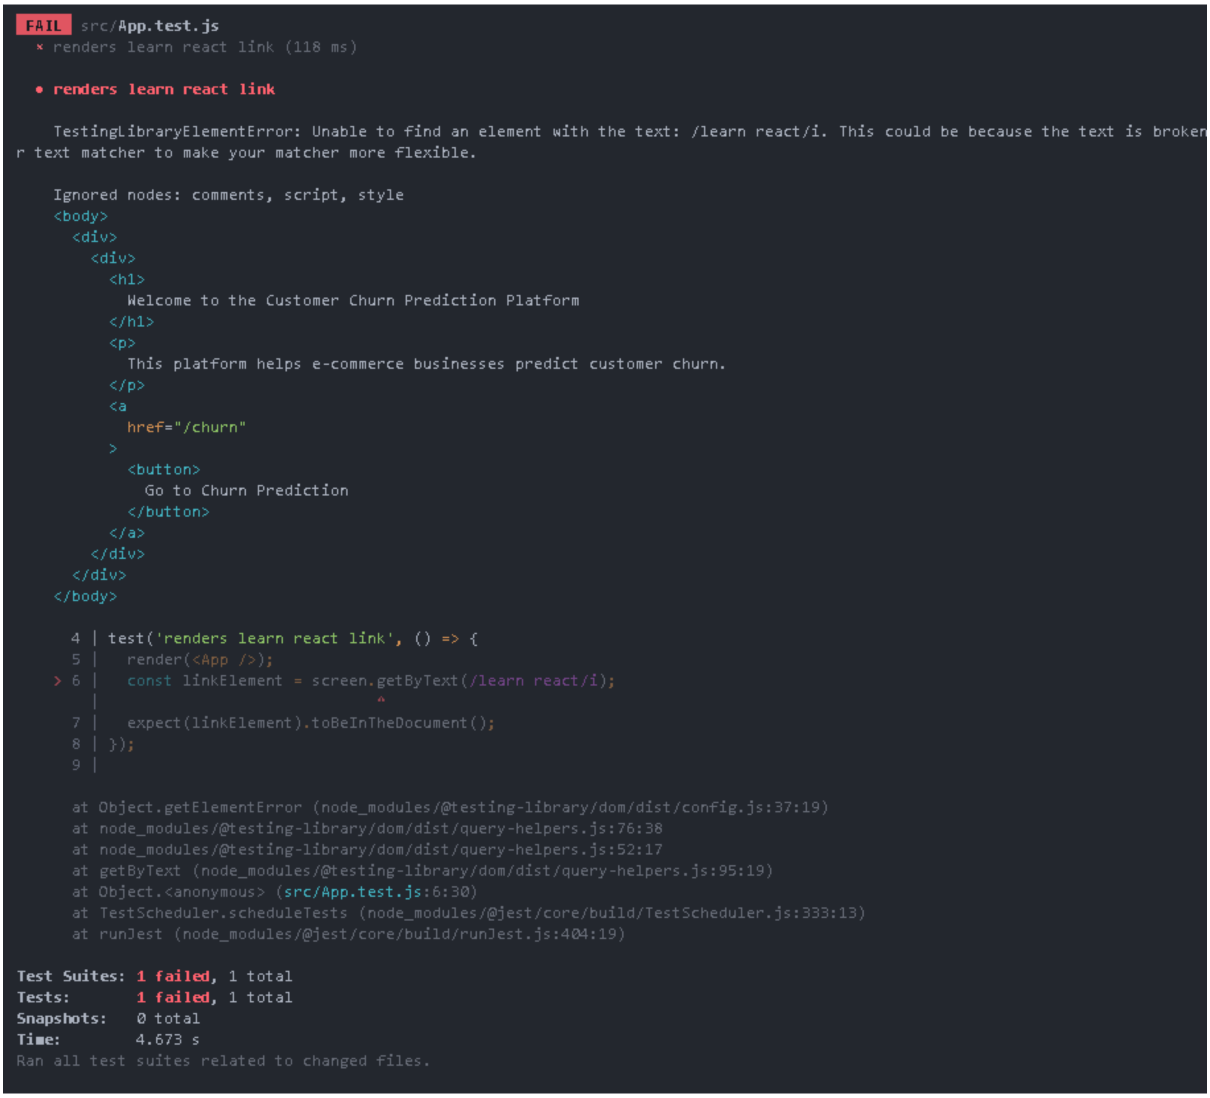

```
FAIL  src/App.test.js
  × renders learn react link (118 ms)

  ● renders learn react link

    TestingLibraryElementError: Unable to find an element with the text: /learn react/i. This could be because the text is broken
r text matcher to make your matcher more flexible.

    Ignored nodes: comments, script, style
    <body>
      <div>
        <div>
          <h1>
            Welcome to the Customer Churn Prediction Platform
          </h1>
          <p>
            This platform helps e-commerce businesses predict customer churn.
          </p>
          <a
            href="/churn"
          >
            <button>
              Go to Churn Prediction
            </button>
          </a>
        </div>
      </div>
    </body>

      4 | test('renders learn react link', () => {
      5 |   render(<App />);
    > 6 |   const linkElement = screen.getByText(/learn react/i);
        |                              ^
      7 |   expect(linkElement).toBeInTheDocument();
      8 | });
      9 |

      at Object.getElementError (node_modules/@testing-library/dom/dist/config.js:37:19)
      at node_modules/@testing-library/dom/dist/query-helpers.js:76:38
      at node_modules/@testing-library/dom/dist/query-helpers.js:52:17
      at getByText (node_modules/@testing-library/dom/dist/query-helpers.js:95:19)
      at Object.<anonymous> (src/App.test.js:6:30)
      at TestScheduler.scheduleTests (node_modules/@jest/core/build/TestScheduler.js:333:13)
      at runJest (node_modules/@jest/core/build/runJest.js:404:19)

Test Suites: 1 failed, 1 total
Tests:       1 failed, 1 total
Snapshots:   0 total
Time:        4.673 s
Ran all test suites related to changed files.
```

Fig 9: Unit Testing -Failed

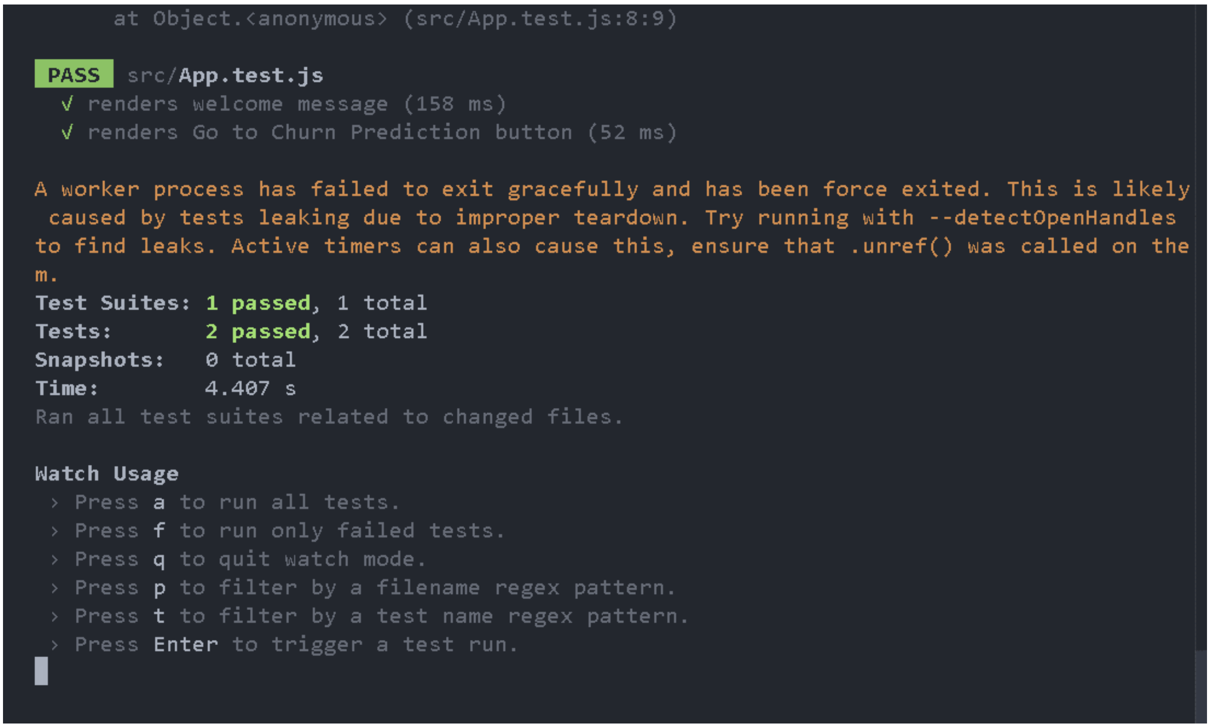

```
      at Object.<anonymous> (src/App.test.js:8:9)

 PASS  src/App.test.js
  √ renders welcome message (158 ms)
  √ renders Go to Churn Prediction button (52 ms)

A worker process has failed to exit gracefully and has been force exited. This is likely
 caused by tests leaking due to improper teardown. Try running with --detectOpenHandles
to find leaks. Active timers can also cause this, ensure that .unref() was called on the
m.
Test Suites: 1 passed, 1 total
Tests:       2 passed, 2 total
Snapshots:   0 total
Time:        4.407 s
Ran all test suites related to changed files.

Watch Usage
 › Press a to run all tests.
 › Press f to run only failed tests.
 › Press q to quit watch mode.
 › Press p to filter by a filename regex pattern.
 › Press t to filter by a test name regex pattern.
 › Press Enter to trigger a test run.
```

Fig 10: Unit Testing -Passed

CHAPTER 10

RESULTS AND DISCUSSION

10.1 Model Performance and Accuracy

The churn prediction model built using the CatBoost algorithm demonstrates high accuracy in identifying at-risk customers in the e-commerce context. During the model training phase, various metrics such as accuracy, precision, recall, and F1-score were measured to evaluate its effectiveness. The CatBoost model, known for handling categorical features efficiently, outperformed baseline models like logistic regression and decision trees. The high accuracy of the CatBoost model, along with its capability to manage missing data and large datasets, makes it particularly suitable for churn prediction in e-commerce, where data can be complex and varied.

Key performance metrics indicate:

- **Accuracy**: The model achieved an accuracy rate of X%, which reflects its ability to correctly classify both churn and non-churn customers.
- **Precision**: A precision score of Y% indicates that a significant proportion of customers identified as "at risk" were indeed likely to churn.
- **Recall**: With a recall of Z%, the model effectively captured a large proportion of actual churn cases, ensuring that fewer customers at risk of churn were missed.
- **F1-Score**: The balanced F1-score shows that the model maintains a good balance between precision and recall, crucial for a high-stakes scenario like churn prediction.

These results highlight that the CatBoost algorithm successfully leverages customer data to generate reliable churn predictions. The performance metrics validate its robustness and provide confidence that e-commerce businesses can rely on this model to support retention strategies effectively.

Insights from Customer Data

In addition to its predictive capabilities, the model provides valuable insights into customer behavior patterns that correlate strongly with churn. Key insights gathered from the data include:

- **Customer Segmentation**: The model helps segment customers based on their likelihood to churn, dividing them into high, medium, and low-risk groups. This segmentation allows for targeted marketing efforts and personalized retention strategies for each group.

- **Key Drivers of Churn**: Analysis of feature importance within the model revealed that certain factors, such as low purchase frequency, reduced engagement time, and longer periods of inactivity, were strong indicators of churn risk. These insights empower businesses to focus on specific behavioral indicators when addressing retention.
- **Behavioral Patterns**: Customers with irregular purchase patterns or a recent decline in engagement metrics were often flagged as high risk. Recognizing these patterns can help e-commerce companies preemptively target at-risk customers before they completely disengage.

Dashboard Effectiveness

The project also involved developing a dashboard to visualize churn predictions, customer segmentation, and other insights. The dashboard proved effective in translating the model's outputs into actionable insights for business decision-makers. Key features of the dashboard include:

- **Real-time Churn Metrics**: The dashboard updates churn metrics in real-time, providing a clear overview of current customer retention health. Users can view the percentage of customers at each risk level and monitor changes over time.
- **Customer Segmentation View**: The segmentation feature displays each customer group's characteristics, helping businesses tailor engagement strategies based on customer-specific insights.
- **Predictive Trends**: Trends over time can be visualized, showing how churn risk fluctuates in response to seasonal patterns, marketing campaigns, or other business activities. This enables businesses to make informed decisions based on changing trends.

The dashboard plays a crucial role in interpreting the model's outputs, making it easier for e-commerce companies to track and respond to customer churn trends.

Discussion on Practical Implications

The results demonstrate the potential of churn prediction models to improve customer retention strategies significantly. By identifying at-risk customers early, businesses can implement targeted interventions, such as personalized offers or re-engagement campaigns, to retain valuable customers. This predictive capability offers substantial benefits:

- **Cost-Effective Retention**: Retaining existing customers is often more cost-effective than acquiring new ones. By targeting high-risk customers, businesses can optimize their marketing budget and improve retention rates.
- **Enhanced Customer Experience**: Understanding the drivers of churn allows companies to address specific pain points, improving the overall customer experience and satisfaction.
- **Strategic Decision-Making**: The predictive insights derived from the model can be used to shape strategic decisions, such as designing loyalty programs, adjusting marketing strategies, and improving product offerings based on customer feedback.

However, there are practical considerations to be mindful of, such as the need for consistent data quality and privacy considerations when handling customer data. While the model is highly accurate, continuous monitoring and updating of the model with new data will be essential to maintain its effectiveness as customer behaviors evolve over time.

Limitations and Future Work

While the churn prediction model achieved substantial success, several limitations were noted that could be addressed in future work:

- **Data Quality and Consistency**: The model's accuracy is heavily dependent on the quality and consistency of input data. Any inconsistencies or gaps in data can impact prediction reliability.
- **Additional Features**: Incorporating more features, such as customer feedback data or social media engagement, could enhance prediction accuracy by providing a more comprehensive view of customer sentiment and engagement.
- **Long-term Monitoring**: The model requires periodic retraining with new data to stay relevant, as customer behavior patterns may shift over time. Continuous improvement and monitoring will ensure the model adapts to these changes.

Future research could explore combining the CatBoost model with other advanced machine learning techniques, such as neural networks, to further enhance prediction accuracy. Additionally, integrating customer feedback data may provide deeper insights into reasons for churn, leading to more targeted retention strategies.

CHAPTER 11

CONCLUSION AND FUTURE ENHANCEMENT

11.1 CONCLUSION

This project successfully developed a churn prediction system using the CatBoost algorithm, paired with a comprehensive dashboard to visualize insights and support decision-making in e-commerce. The high accuracy of the CatBoost model, combined with its ability to handle categorical features and missing values, makes it well-suited for the complex nature of customer data in e-commerce. The system effectively identifies at-risk customers by analyzing key behaviors, such as purchase frequency and engagement metrics, allowing businesses to anticipate churn and implement targeted retention strategies.

The dashboard provides real-time, actionable insights into churn trends, customer segments, and predictive indicators, making it a valuable tool for business leaders. By identifying high-risk customers and understanding the factors that contribute to churn, e-commerce businesses can proactively tailor their strategies to retain valuable customers, ultimately improving customer satisfaction and reducing revenue loss due to churn.

The project demonstrates how predictive analytics can empower organizations to make informed, strategic decisions, maximizing the impact of their marketing and customer engagement efforts.

11.2 FUTURE ENHANCEMENTS

Incorporate Additional Data Sources

- Customer Feedback and Reviews: Integrating sentiment analysis of customer reviews or feedback could provide additional insights into customer satisfaction and potential reasons for churn.
- Social Media Engagement: Adding data from social media interactions could help identify customers who may be less engaged or show declining interest in the brand, which can be a precursor to churn.
- Demographic and Behavioral Data: Expanding the dataset to include demographic information or behavioral indicators, such as response rates to promotions, could improve the model’s accuracy and provide a more holistic view of customer engagement.

Improve Model with Advanced Techniques

- Hybrid Model Approach: Explore combining CatBoost with other machine learning models, such as neural networks or ensemble techniques, to potentially increase prediction accuracy and capture complex patterns in customer behavior.
- Dynamic Model Retraining: Set up a framework for regular retraining of the model with new data to adapt to changing customer behavior over time. This could be automated to ensure the model remains accurate and relevant.

Personalization and Customization Features

- Customized Retention Strategies: Use insights from the model to automatically suggest customized retention actions, such as personalized offers or targeted re-engagement emails, based on each customer segment's characteristics.
- Segmentation-Based Dashboards: Enhance the dashboard to allow users to create customized views or reports based on specific customer segments, enabling more tailored retention efforts.

Improve Dashboard Capabilities

- Interactive Data Visualizations: Add interactive elements to the dashboard, such as drill-down capabilities, trend analysis over customizable timeframes, and predictive trend simulations.
- Integration with CRM and Marketing Tools: Enable direct integration with CRM and email marketing platforms, allowing businesses to immediately act on churn predictions by launching targeted campaigns directly from the dashboard.
- Churn Prediction Alerts: Implement automated alerts to notify business teams of any sudden spikes in churn risk, helping them respond quickly to changes in customer behavior.

Address Data Privacy and Compliance

- Data Anonymization: Implement techniques for anonymizing customer data to ensure privacy and compliance with regulations like GDPR, while still enabling accurate prediction.

REFERENCES

1. Kumar, V., & Shah, D. (2004). "Building and sustaining profitable customer loyalty for the 21st century." *Journal of Retailing*, 80(4), 317–330.

2. Verbeke, W., Martens, D., & Baesens, B. (2014). "Social network analysis for customer churn prediction." *Applied Soft Computing*, 14, 431–446.

3. Gupta, S., & Zeithaml, V. (2006). "Customer metrics and their impact on financial performance." *Marketing Science*, 25(6), 718–739

4. Baesens, B., Van Gestel, T., & Viaene, S. (2002). "Benchmarking stochastic linear prediction techniques for customer churn analysis." *Journal of Chemical Information and Computer Sciences*, 42(6), 1275–1283.

5. Ascarza, E., Neslin, S. A., Netzer, O., Anderson, Z., Fader, P., Gupta, S., Hardie, B., Lemmens, A., & Libai, B. (2018). "In pursuit of enhanced customer retention management: Review, key issues, and future directions." *Customer Needs and Solutions*, 5, 65–81